

Clustering Problem

- Find similar countries based on different attributes like Income per capita, phones per 1000, GDP, etc. and create clusters for the same. Plot the cluster on the world map using the ISO codes provided, depicting the clusters created.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df= pd.read_csv("/content/drive/MyDrive/MSBA/Semester 1/IDS-572-DataMiningProfNegar/project winter break/Country_Facts.csv")
df.head()
```

	Country	Region	Population	Area (sq. mi.)	Pop. Density (per sq. mi.)	Coastline (coast/area ratio)	Net migration	Infant mortality (per 1000 births)	GDP (\$ per capita)	Literacy (%)	Phones (per 1000)	Arable (%)
0	Afghanistan	ASIA (EX. NEAR EAST)	31056997	647500	48.0	0.00	23.06	163.07	700.0	36.0	3.2	12.1
1	Albania	EASTERN EUROPE	3581655	28748	124.6	1.26	-4.93	21.52	4500.0	86.5	71.2	21.0
2	Algeria	NORTHERN AFRICA	32930091	2381740	13.8	0.04	-0.39	31.00	6000.0	70.0	78.1	3.2
3	American Samoa	OCEANIA	57794	199	290.4	58.29	-20.71	9.27	8000.0	97.0	259.5	10.0
4	Andorra	WESTERN EUROPE	71201	468	152.1	0.00	6.60	4.05	19000.0	100.0	497.2	2.2



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 227 entries, 0 to 226
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               227 non-null    object
1   Region                                227 non-null    object
2   Population                             227 non-null    int64
3   Area (sq. mi.)                         227 non-null    int64
```

4	Pop. Density (per sq. mi.)	227 non-null	float64
5	Coastline (coast/area ratio)	227 non-null	float64
6	Net migration	224 non-null	float64
7	Infant mortality (per 1000 births)	224 non-null	float64
8	GDP (\$ per capita)	226 non-null	float64
9	Literacy (%)	209 non-null	float64
10	Phones (per 1000)	223 non-null	float64
11	Arable (%)	225 non-null	float64
12	Crops (%)	225 non-null	float64
13	Other (%)	225 non-null	float64
14	Climate	205 non-null	float64
15	Birthrate	224 non-null	float64
16	Deathrate	223 non-null	float64
17	Agriculture	212 non-null	float64
18	Industry	211 non-null	float64
19	Service	212 non-null	float64

dtypes: float64(16), int64(2), object(2)

memory usage: 35.6+ KB

df.columns

```
Index(['Country', 'Region', 'Population', 'Area (sq. mi.)',
      'Pop. Density (per sq. mi.)', 'Coastline (coast/area ratio)',
      'Net migration', 'Infant mortality (per 1000 births)',
      'GDP ($ per capita)', 'Literacy (%)', 'Phones (per 1000)', 'Arable (%)',
      'Crops (%)', 'Other (%)', 'Climate', 'Birthrate', 'Deathrate',
      'Agriculture', 'Industry', 'Service'],
      dtype='object')
```

Double-click (or enter) to edit

df.shape

(227, 20)

len(df['Country'].unique())


227

df.describe().transpose()

	count	mean	std	min	25%	50%	75%	max
Population	227.0	2.874028e+07	1.178913e+08	7026.000	437624.00000	4786994.000	1.749777e+07	1.313974e+09
Area (sq. mi.)	227.0	5.982270e+05	1.790282e+06	2.000	4647.50000	86600.000	4.418110e+05	1.707520e+07
Pop. Density (per sq. mi.)	227.0	3.790471e+02	1.660186e+03	0.000	29.15000	78.800	1.901500e+02	1.627150e+04
Coastline (coast/area ratio)	227.0	2.116533e+01	7.228686e+01	0.000	0.10000	0.730	1.034500e+01	8.706600e+02
Net migration	224.0	3.812500e-02	4.889269e+00	-20.990	-0.92750	0.000	9.975000e-01	2.306000e+01
Infant mortality (per 1000 births)	224.0	3.550696e+01	3.538990e+01	2.290	8.15000	21.000	5.570500e+01	1.911900e+02
GDP (\$ per capita)	226.0	9.689823e+03	1.004914e+04	500.000	1900.00000	5550.000	1.570000e+04	5.510000e+04
Literacy (%)	209.0	8.283828e+01	1.972217e+01	17.600	70.60000	92.500	9.800000e+01	1.000000e+02
Phones (per 1000)	223.0	2.360614e+02	2.279918e+02	0.200	37.80000	176.200	3.896500e+02	1.035600e+03

```
sns.displot(data=df,x= 'Population')
```

```
<seaborn.axisgrid.FacetGrid at 0x7fc41d7d2d30>
```



Population is data is very skewed and not normally distributed


```
df['Region'].value_counts()
```

```
SUB-SAHARAN AFRICA      51
LATIN AMER. & CARIB     45
ASIA (EX. NEAR EAST)    28
WESTERN EUROPE          28
OCEANIA                 21
NEAR EAST               16
EASTERN EUROPE          12
C.W. OF IND. STATES     12
NORTHERN AFRICA         6
NORTHERN AMERICA        5
BALTICS                 3
```

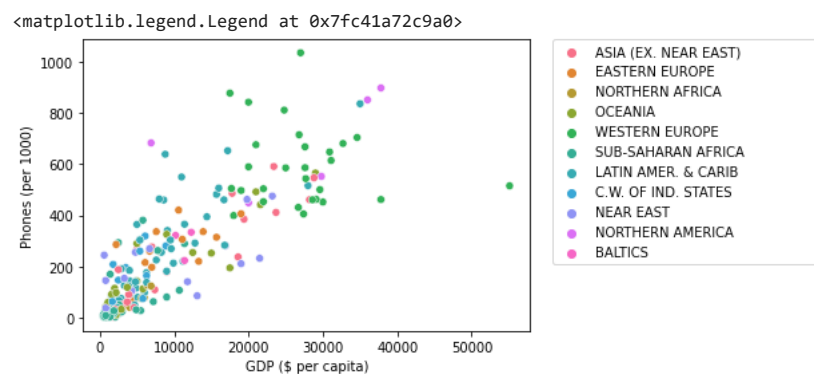
```
Name: Region, dtype: int64
```

```
plt.figure(figsize=(10,6),dpi=150)
sns.barplot(data=df,y='GDP ($ per capita)',x='Region',estimator=np.mean)
plt.xticks(rotation=90)

(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
<a list of 11 Text major ticklabel objects>)
```



```
sns.scatterplot(data=df,x= 'GDP ($ per capita)',y= 'Phones (per 1000)',hue='Region')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0)
```

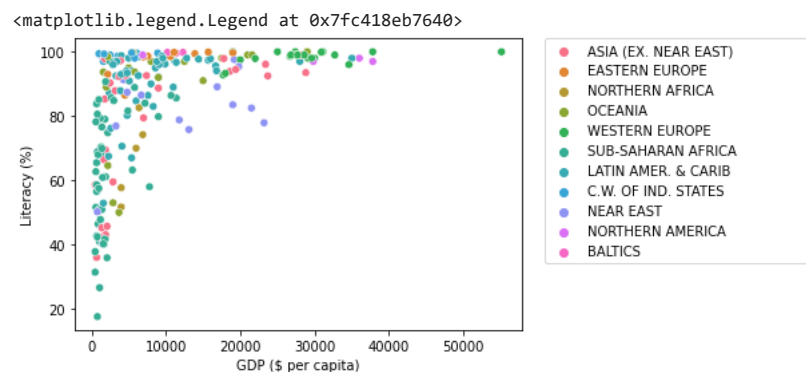


```
df[(df['GDP ($ per capita)'] > 30000) & (df['Phones (per 1000)'] < 600)][['Country']]
```

```
121    Luxembourg
154    Norway
Name: Country, dtype: object
```

These 2 countries have high GDP but less Number of Phones. An interesting finding

```
sns.scatterplot(data=df,x='GDP ($ per capita)',y='Literacy (%)',hue='Region')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0)
```




```
df = df.drop('Other (%)',axis=1)
```

```
df.isnull().sum()
```

```
Country      0
Region       0
Population   0
```

```
Area (sq. mi.)          0
Pop. Density (per sq. mi.) 0
Coastline (coast/area ratio) 0
Net migration           3
Infant mortality (per 1000 births) 3
GDP ($ per capita)      1
Literacy (%)            18
Phones (per 1000)       4
Arable (%)              2
Crops (%)               2
Climate                 22
Birthrate               3
Deathrate               4
Agriculture             15
Industry                16
Service                 15
dtype: int64
```

```
sns.heatmap(df.corr(), cmap='Paired')
```



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 227 entries, 0 to 226
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               227 non-null    object
1   Region                               227 non-null    object
2   Population                            227 non-null    int64
3   Area (sq. mi.)                        227 non-null    int64
4   Pop. Density (per sq. mi.)            227 non-null    float64
5   Coastline (coast/area ratio)           227 non-null    float64
6   Net migration                          224 non-null    float64
7   Infant mortality (per 1000 births)      224 non-null    float64
8   GDP ($ per capita)                     226 non-null    float64
```

```
9 Literacy (%) 209 non-null float64
10 Phones (per 1000) 223 non-null float64
11 Arable (%) 225 non-null float64
12 Crops (%) 225 non-null float64
13 Climate 205 non-null float64
14 Birthrate 224 non-null float64
15 Deathrate 223 non-null float64
16 Agriculture 212 non-null float64
17 Industry 211 non-null float64
18 Service 212 non-null float64
dtypes: float64(15), int64(2), object(2)
memory usage: 33.8+ KB
```

```
df[df['Agriculture'].isnull()][['Country','GDP ($ per capita)']]
```

	Country	GDP (\$ per capita)	
3	American Samoa	8000.0	
4	Andorra	19000.0	
78	Gibraltar	17500.0	
80	Greenland	20000.0	
83	Guam	21000.0	
134	Mayotte	2600.0	
140	Montserrat	3400.0	
144	Nauru	5000.0	
153	N. Mariana Islands	12500.0	
171	Saint Helena	2500.0	
174	St Pierre & Miquelon	6900.0	
177	San Marino	34600.0	
208	Turks & Caicos Is	9600.0	
221	Wallis and Futuna	3700.0	
223	Western Sahara	NaN	

```
df[df['Industry'].isnull()][['Country','GDP ($ per capita)']]
```

	Country	GDP (\$ per capita)
3	American Samoa	8000.0
4	Andorra	19000.0
78	Gibraltar	17500.0
80	Greenland	20000.0
83	Guam	21000.0
134	Mayotte	2600.0
138	Monaco	27000.0
140	Montserrat	3400.0

```
df[df['Agriculture'].isnull()] = df[df['Agriculture'].isnull()].fillna(0)
```

153	N. Mariana Islands	12500.0
-----	--------------------	---------

```
df.isnull().sum()
```

Country	0
Region	0
Population	0
Area (sq. mi.)	0
Pop. Density (per sq. mi.)	0
Coastline (coast/area ratio)	0
Net migration	1
Infant mortality (per 1000 births)	1
GDP (\$ per capita)	0
Literacy (%)	13
Phones (per 1000)	2
Arable (%)	1
Crops (%)	1
Climate	18
Birthrate	1
Deathrate	2
Agriculture	0
Industry	1
Service	1
dtype: int64	

```
df['Climate'] = df['Climate'].fillna(df.groupby('Region')['Climate'].transform('mean'))
```

```
df.isnull().sum()
```

Country	0
Region	0
Population	0
Area (sq. mi.)	0
Pop. Density (per sq. mi.)	0
Coastline (coast/area ratio)	0
Net migration	1
Infant mortality (per 1000 births)	1
GDP (\$ per capita)	0
Literacy (%)	13
Phones (per 1000)	2
Arable (%)	1
Crops (%)	1
Climate	0


```

Birthrate      1
Deathrate      2
Agriculture    0
Industry        1
Service         1
dtype: int64

```

```
df['Literacy (%)'] = df['Literacy (%)'].fillna(df.groupby('Region')['Literacy (%)'].transform('mean'))
```

```
df.isnull().sum()
```

```

Country      0
Region       0
Population   0
Area (sq. mi.) 0
Pop. Density (per sq. mi.) 0
Coastline (coast/area ratio) 0
Net migration 1
Infant mortality (per 1000 births) 1
GDP ($ per capita) 0
Literacy (%)  0
Phones (per 1000) 2
Arable (%)    1
Crops (%)     1
Climate       0
Birthrate     1
Deathrate     2
Agriculture   0
Industry      1
Service       1
dtype: int64

```

```
df= df.dropna()
```

```
df.isnull().sum()
```

```

Country      0
Region       0
Population   0
Area (sq. mi.) 0
Pop. Density (per sq. mi.) 0
Coastline (coast/area ratio) 0
Net migration 0
Infant mortality (per 1000 births) 0
GDP ($ per capita) 0
Literacy (%)  0
Phones (per 1000) 0
Arable (%)    0
Crops (%)     0
Climate       0
Birthrate     0
Deathrate     0
Agriculture   0
Industry      0
Service       0
dtype: int64

```

All missing values handled!

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 221 entries, 0 to 226
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               221 non-null    object
1   Region                               221 non-null    object
2   Population                            221 non-null    int64
3   Area (sq. mi.)                       221 non-null    int64
4   Pop. Density (per sq. mi.)           221 non-null    float64
5   Coastline (coast/area ratio)         221 non-null    float64
6   Net migration                        221 non-null    float64
7   Infant mortality (per 1000 births)   221 non-null    float64
8   GDP ($ per capita)                   221 non-null    float64
9   Literacy (%)                         221 non-null    float64
10  Phones (per 1000)                    221 non-null    float64
11  Arable (%)                           221 non-null    float64
12  Crops (%)                            221 non-null    float64
13  Climate                              221 non-null    float64
14  Birthrate                            221 non-null    float64
15  Deathrate                            221 non-null    float64
16  Agriculture                           221 non-null    float64
17  Industry                             221 non-null    float64
18  Service                              221 non-null    float64
dtypes: float64(15), int64(2), object(2)
memory usage: 34.5+ KB
```

```
X = df.drop('Country',axis=1)
```

```
X = pd.get_dummies(X)
```

```
X.head()
```

	Population	Area (sq. mi.)	Pop. Density (per sq. mi.)	Coastline (coast/area ratio)	Net migration	Infant mortality (per 1000 births)	GDP (\$ per capita)	Literacy (%)	Phones (per 1000)	Arable (%)	...	Region_BALTICS	Region_C.W. OF IND. STATES	Region_EASTERN EUROPE	Region_LATIN AMER. & CARIB	Region_NEAR EAST	Regi
0	31056997	647500	48.0	0.00	23.06	163.07	700.0	36.0	3.2	12.13	...	0	0	0	0	0	
1	3581655	28748	124.6	1.26	-4.93	21.52	4500.0	86.5	71.2	21.09	...	0	0	1	0	0	
2	32930091	2381740	13.8	0.04	-0.39	31.00	6000.0	70.0	78.1	3.22	...	0	0	0	0	0	
3	57794	199	290.4	58.29	-20.71	9.27	8000.0	97.0	259.5	10.00	...	0	0	0	0	0	
4	71201	468	152.1	0.00	6.60	4.05	19000.0	100.0	497.2	2.22	...	0	0	0	0	0	

5 rows × 28 columns



Scaling the features of X dataframe

```

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

scaled_X = scaler.fit_transform(X)

scaled_X

array([[2.36307140e-02, 3.79200985e-02, 2.96607551e-03, ...,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
       [2.72048678e-03, 1.68320206e-03, 7.69943768e-03, ...,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
       [2.50562403e-02, 1.39484983e-01, 8.52746710e-04, ...,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
       ...,
       [1.63239770e-02, 3.09198848e-02, 2.50880554e-03, ...,
        0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
       [8.74830703e-03, 4.40760465e-02, 9.45436569e-04, ...,
        0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
       [9.30752592e-03, 2.28737092e-02, 1.93412841e-03, ...,
        0.00000000e+00, 1.00000000e+00, 0.00000000e+00]])

from sklearn.cluster import KMeans

inertias = []

for i in range(2,30):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(scaled_X)
    inertias.append(kmeans.inertia_)

plt.plot(range(2,30), inertias, marker='o')
plt.title('Elbow method')
plt.xlabel('Number of clusters')
plt.ylabel('Sum of Squared Distances')
plt.show()

```

```
... ..
kmeans = KMeans(n_clusters=15)
kmeans.fit(scaled_X)

KMeans(n_clusters=15)
|
kmeans.labels_

array([[13, 10, 9, 2, 4, 1, 0, 0, 12, 3, 0, 2, 4, 3, 0, 5, 13,
        0, 3, 4, 12, 11, 8, 13, 12, 10, 6, 12, 0, 13, 10, 1, 13, 11,
        13, 11, 8, 6, 0, 1, 1, 12, 13, 12, 11, 1, 6, 12, 11, 10, 12,
        10, 4, 6, 12, 12, 12, 9, 12, 6, 6, 14, 1, 4, 2, 4, 4, 0,
        2, 6, 11, 5, 3, 4, 11, 4, 4, 8, 12, 0, 2, 12, 1, 1, 12,
        12, 12, 7, 10, 4, 13, 13, 13, 5, 4, 4, 5, 4, 12, 7, 4, 5,
        3, 6, 2, 13, 7, 5, 3, 13, 14, 5, 6, 1, 9, 4, 14, 4, 7,
        10, 11, 11, 13, 13, 1, 4, 2, 0, 1, 11, 1, 12, 2, 3, 13, 12,
        9, 1, 6, 2, 13, 4, 0, 2, 2, 12, 1, 11, 2, 4, 5, 13, 2,
        12, 2, 12, 12, 13, 10, 4, 12, 5, 6, 10, 3, 11, 6, 0, 12, 8,
        12, 2, 4, 11, 5, 11, 6, 1, 7, 10, 10, 2, 1, 6, 4, 13, 11,
        12, 6, 4, 4, 5, 7, 3, 1, 13, 11, 2, 12, 9, 5, 3, 12, 2,
        11, 3, 5, 4, 8, 12, 3, 2, 12, 13, 0, 2, 5, 9, 5, 6, 6],
      dtype=int32)
```

```
iso_codes = pd.read_csv('/content/drive/MyDrive/MSBA/Semester 1/IDS-572-DataMiningProfNegar/project winter break/country_iso_codes.csv')
```

```
iso_codes.head()
```

	Country	ISO Code
0	Afghanistan	AFG
1	Akrotiri and Dhekelia – See United Kingdom, The	Akrotiri and Dhekelia – See United Kingdom, The
2	Åland Islands	ALA
3	Albania	ALB
4	Algeria	DZA

```
iso_mapping = iso_codes.set_index('Country')
iso_mapping.head()
```

	Country	ISO Code
	Afghanistan	AFG
	Akrotiri and Dhekelia – See United Kingdom, The	Akrotiri and Dhekelia – See United Kingdom, The
	Åland Islands	ALA
	Albania	ALB
	Algeria	DZA

```
iso_mapping_dict = iso_mapping['ISO Code'].to_dict()
```

```
iso_mapping_dict
```

```
'Switzerland': 'CHE',
'Syrian Arab Republic (the)\u200a[x]': 'SYR',
'Taiwan (Province of China)\u200a[y]': 'TWN',
'Tajikistan': 'TJK',
'Tanzania, the United Republic of': 'TZA',
'Thailand': 'THA',
'Timor-Leste\u200a[aa]': 'TLS',
'Togo': 'TGO',
'Tokelau': 'TKL',
'Tonga': 'TON',
'Trinidad and Tobago': 'TTO',
'Tunisia': 'TUN',
'Turkey': 'TUR',
'Turkmenistan': 'TKM',
'Turks and Caicos Islands (the)': 'TCA',
'Tuvalu': 'TUV',
'Uganda': 'UGA',
'Ukraine': 'UKR',
'United Arab Emirates (the)': 'ARE',
'United Kingdom of Great Britain and Northern Ireland (the)': 'GBR',
'United States Minor Outlying Islands (the)\u200a[ac]': 'UMI',
'United States of America (the)': 'USA',
'United States Virgin Islands - See Virgin Islands (U.S.)': 'United States Virgin Islands - See Virgin Islands (U.S.)',
'Uruguay': 'URY',
'Uzbekistan': 'UZB',
'Vanuatu': 'VUT',
'Vatican City - See Holy See, The.': 'Vatican City - See Holy See, The.',
'Venezuela (Bolivarian Republic of)': 'VEN',
'Viet Nam\u200a[ae]': 'VNM',
'Virgin Islands (British)\u200a[af]': 'VGB',
'Virgin Islands (U.S.)\u200a[ag]': 'VIR',
'Wales - See United Kingdom, The.': 'Wales - See United Kingdom, The.',
'Wallis and Futuna': 'WLF',
'Western Sahara\u200a[ah]': 'ESH',
'Yemen': 'YEM',
'Zambia': 'ZMB',
'Zimbabwe': 'ZWE',
'United States': 'USA',
'United Kingdom': 'GBR',
'Venezuela': 'VEN',
'Australia': 'AUS',
'Iran': 'IRN',
'France': 'FRA',
'Russia': 'RUS',
'Korea, North': 'PRK',
'Korea, South': 'KOR',
'Myanmar': 'MMR',
'Burma': 'MMR',
'Vietnam': 'VNM',
'Laos': 'LAO',
'Bolivia': 'BOL',
'Niger': 'NER',
'Sudan': 'SDN',
'Congo, Dem. Rep.': 'COD',
'Congo, Repub. of the': 'COG',
'Tanzania': 'TZA',
'Central African Rep.': 'CAF',
'Cote d'Ivoire': 'CIV'}
```

```
df['ISO Code'] = df['Country'].map(iso_mapping_dict)
```

```
df[['Country', 'ISO Code']].head()
```

	Country	ISO Code
0	Afghanistan	AFG
1	Albania	ALB
2	Algeria	DZA
3	American Samoa	ASM
4	Andorra	AND

```
df['Cluster'] = kmeans.labels_
```

```
import plotly.express as px
```

```
fig = px.choropleth(df, locations='ISO Code', color='Cluster', hover_name='Country',
                    color_continuous_scale='Viridis_r')
```

```
fig.update_layout(title_text="Clustering Countries based on K-Means Algorithm")
fig.show()
```

Clustering Countries based on K-Means Algorithm



✓ 1s completed at 9:28 PM

● ✕